

ECE 565 homework 1

Js895

Junqi Sun

Question1 :

(a)

| Exe time(millisecons) | Optimization level | | Original output |
|-----------------------|--------------------|---------|-----------------|
| Cmd line argument | O2 | O3 | No optimization |
| 10,000,000 | 22.885 | 18.366 | 77.973 |
| 100,000,000 | 227.108 | 184.820 | 776.493 |

All result are minimum value from 5 tests

Original O2 assembly code:

```
for (i=N-1; i>=1; i--) {
1474:      8d 41 ff          lea    -0x1(%rcx),%eax
1477:      85 c0            test   %eax,%eax
1479:      7e 11           jle    148c <do_loops+0x1c>
147b:      48 98           cltq
147d:      0f 1f 00        nopl   (%rax)
a[i] = a[i] + 1;
1480:      83 04 87 01     addl   $0x1, (%rdi,%rax,4)
for (i=N-1; i>=1; i--) {
1484:      48 83 e8 01     sub    $0x1,%rax
1488:      85 c0            test   %eax,%eax
148a:      7f f4           jg     1480 <do_loops+0x10>
```

```
for (i=1; i<N; i++) {
148c:      83 f9 01          cmp    $0x1,%ecx
148f:      7e 3d           jle    14ce <do_loops+0x5e>
1491:      44 8d 41 fe     lea    -0x2(%rcx),%r8d
1495:      b8 02 00 00 00    mov    $0x2,%eax
149a:      4d 8d 48 03     lea    0x3(%r8),%r9
149e:      66 90           xchg   %ax,%ax
b[i] = a[i+1] + 3;
14a0:      8b 0c 87         mov    (%rdi,%rax,4),%ecx
14a3:      83 c1 03         add    $0x3,%ecx
14a6:      89 4c 86 fc     mov    %ecx,-0x4(%rsi,%rax,4)
```

```

for (i=1; i<N; i++) {
14aa:    48 83 c0 01    add    $0x1,%rax
14ae:    49 39 c1       cmp    %rax,%r9
14b1:    75 ed         jne    14a0 <do_loops+0x30>
14b3:    31 c0       xor    %eax,%eax
14b5:    0f 1f 00     nopl   (%rax)
c[i] = b[i-1] + 2;
14b8:    8b 3c 86     mov    (%rsi,%rax,4),%edi
14bb:    8d 4f 02     lea    0x2(%rdi),%ecx
14be:    89 4c 82 04   mov    %ecx,0x4(%rdx,%rax,4)
for (i=1; i<N; i++) {
14c2:    48 89 c1     mov    %rax,%rcx
14c5:    48 83 c0 01   add    $0x1,%rax
14c9:    4c 39 c1     cmp    %r8,%rcx
14cc:    75 ea         jne    14b8 <do_loops+0x48>
}

```

(b)

Processor architecture: x86_64

CPU frequency: 2.3GHZ

OS type: Linux

VM: Ubuntu 20.04.1 LTS

(c)

Optimization 1: loop unrolling

Source code:

```

// loop unrolling
int i;
for (i=N-1; i>=1; i-=4) {
    a[i] = a[i] + 1;
    a[i-1] = a[i-1] + 1;
    a[i-2] = a[i-2] + 1;
    a[i-3] = a[i-3] + 1;
}

for (i=1; i<N; i+=4) {
    b[i] = a[i+1] + 3;
    b[i+1] = a[i+2] + 3;
    b[i+2] = a[i+3] + 3;
    b[i+3] = a[i+4] + 3;
}

for (i=1; i<N; i+=4) {
    c[i] = b[i-1] + 2;
    c[i+1] = b[i] + 2;
    c[i+2] = b[i+1] + 2;
    c[i+3] = b[i+2] + 2;
}

```

Performance result:

| Exe time(milliseconds) | Optimization level | |
|------------------------|--------------------|---------|
| Cmd line argument | O2 | O3 |
| 10,000,000 | 21.402 | 19.751 |
| 100,000,000 | 208.817 | 199.864 |

Assembly Code of O2:

```

for (i=N-1; i>=1; i-=4) {
1474:      8d 41 ff      lea    -0x1(%rcx),%eax
1477:      85 c0        test   %eax,%eax
1479:      7e 3d        jle    14b8 <do_loops+0x48>
147b:      48 98        cltq
147d:      44 8d 49 fe    lea    -0x2(%rcx),%r9d
1481:      4c 8d 04 85 00 00 00 lea    0x0(,%rax,4),%r8
1488:      00
1489:      41 c1 e9 02     shr    $0x2,%r9d
148d:      4a 8d 04 07     lea    (%rdi,%r8,1),%rax
1491:      49 c1 e1 04     shl    $0x4,%r9
1495:      4e 8d 44 07 f0 lea    -0x10(%rdi,%r8,1),%r8
149a:      4d 29 c8        sub    %r9,%r8
149d:      0f 1f 00        nopl   (%rax)
a[i] = a[i] + 1;
14a0:      83 00 01        addl   $0x1,(%rax)
a[i-1] = a[i-1] + 1;
14a3:      83 40 fc 01     addl   $0x1,-0x4(%rax)
a[i-2] = a[i-2] + 1;
14a7:      83 40 f8 01     addl   $0x1,-0x8(%rax)
a[i-3] = a[i-3] + 1;
14ab:      83 40 f4 01     addl   $0x1,-0xc(%rax)
for (i=N-1; i>=1; i-=4) {
14af:      48 83 e8 10     sub    $0x10,%rax
14b3:      4c 39 c0        cmp    %r8,%rax
14b6:      75 e8        jne    14a0 <do_loops+0x30>
for (i=1; i<N; i+=4) {
14b8:      83 f9 01        cmp    $0x1,%ecx
14bb:      0f 8e 8f 00 00 00 jle    1550 <do_loops+0xe0>
14c1:      44 8d 49 fe    lea    -0x2(%rcx),%r9d
14c5:      48 8d 47 08     lea    0x8(%rdi),%rax
14c9:      44 89 c9        mov    %r9d,%ecx
14cc:      4c 8d 46 04     lea    0x4(%rsi),%r8
14d0:      c1 e9 02        shr    $0x2,%ecx
14d3:      41 89 c9        mov    %ecx,%r9d
14d6:      49 c1 e1 04     shl    $0x4,%r9
14da:      4a 8d 7c 0f 18 lea    0x18(%rdi,%r9,1),%rdi
14df:      90            nop
b[i] = a[i+1] + 3;
14e0:      8b 08        mov    (%rax),%ecx
14e2:      48 83 c0 10     add    $0x10,%rax
14e6:      49 83 c0 10     add    $0x10,%r8
14ea:      83 c1 03        add    $0x3,%ecx
14ed:      41 89 48 f0     mov    %ecx,-0x10(%r8)
b[i+1] = a[i+2] + 3;
14f1:      8b 48 f4        mov    -0xc(%rax),%ecx
14f4:      83 c1 03        add    $0x3,%ecx
14f7:      41 89 48 f4     mov    %ecx,-0xc(%r8)
b[i+2] = a[i+3] + 3;
14fb:      8b 48 f8        mov    -0x8(%rax),%ecx
14fe:      83 c1 03        add    $0x3,%ecx
1501:      41 89 48 f8     mov    %ecx,-0x8(%r8)
b[i+3] = a[i+4] + 3;
1505:      8b 48 fc        mov    -0x4(%rax),%ecx
1508:      83 c1 03        add    $0x3,%ecx
150b:      41 89 48 fc     mov    %ecx,-0x4(%r8)

```

```

for (i=1; i<N; i+=4) {
150f: 48 39 f8          cmp    %rdi,%rax
1512: 75 cc            jne    14e0 <do_loops+0x70>
1514: 48 83 c2 04      add    $0x4,%rdx
1518: 4a 8d 4c 0e 10   lea    0x10(%rsi,%r9,1),%rcx
151d: 0f 1f 00        nopl   (%rax)
c[i] = b[i-1] + 2;
1520: 8b 06           mov    (%rsi),%eax
1522: 48 83 c6 10     add    $0x10,%rsi
1526: 48 83 c2 10     add    $0x10,%rdx
152a: 83 c0 02       add    $0x2,%eax
152d: 89 42 f0       mov    %eax,-0x10(%rdx)
c[i+1] = b[i] + 2;
1530: 8b 46 f4       mov    -0xc(%rsi),%eax
1533: 83 c0 02       add    $0x2,%eax
1536: 89 42 f4       mov    %eax,-0xc(%rdx)
c[i+2] = b[i+1] + 2;
1539: 8b 46 f8       mov    -0x8(%rsi),%eax
153c: 83 c0 02       add    $0x2,%eax
153f: 89 42 f8       mov    %eax,-0x8(%rdx)
c[i+3] = b[i+2] + 2;
1542: 8b 46 fc       mov    -0x4(%rsi),%eax
1545: 83 c0 02       add    $0x2,%eax
1548: 89 42 fc       mov    %eax,-0x4(%rdx)
for (i=1; i<N; i+=4) {
154b: 48 39 ce       cmp    %rcx,%rsi
154e: 75 d0         jne    1520 <do_loops+0xb0>
}

```

Analysis:

Compare to the original code with O2, O3, the unrolling version of O2 is faster than original while O3 is a bit slower.

When referring to assembly code: In the first loop in **original code**, we use %rax as index i, and do **addl \$0x1, (%rdi, %rax, 4)** every time in loop. While in **unrolling loop**, this step changes to **addl \$0x1, (%rax), addl \$0x1, -0x4(%rax), addl \$0x1, -0x8(%rax), addl \$0x1, -0xc(%rax)**. I think this adjustment here means we don't need to fetch for %rdi every time, we just need to add instant number which is faster.

In order to achieve that, compiler did more preparation step when first enter the loop, which load array address information(%rdi) into %rax before the first operation, thus save the time. The following loop perform similar methods to optimize.

Loop unrolling also expose additional ILP, and reduce instruction count to improve performance.

Optimization 2: loop fusion

Source code:

```
int i;
for (i=N-1; i>=1; i--) {
    a[i] = a[i] + 1;
}
for (i=1; i<N; i++) {
    b[i] = a[i+1] + 3;
    c[i] = b[i-1] + 2;
}
```

Performance result:

| Exe time(milliseconds) | Optimization level | |
|------------------------|--------------------|---------|
| | O2 | O3 |
| Cmd line argument | O2 | O3 |
| 10,000,000 | 22.631 | 22.200 |
| 100,000,000 | 213.774 | 222.919 |

Assembly Code of O2:

```
1470:    f3 0f 1e fa                endbr64
for (i=N-1; i>=1; i--) {
1474:    8d 41 ff                    lea    -0x1(%rcx),%eax
1477:    85 c0                        test   %eax,%eax
1479:    7e 11                        jle    148c <do_loops+0x1c>
147b:    48 98                        cltq
147d:    0f 1f 00                    nopl   (%rax)
a[i] = a[i] + 1;
1480:    83 04 87 01                addl   $0x1,(%rdi,%rax,4)
for (i=N-1; i>=1; i--) {
1484:    48 83 e8 01                sub    $0x1,%rax
1488:    85 c0                        test   %eax,%eax
148a:    7f f4                        jg     1480 <do_loops+0x10>
for (i=1; i<N; i++) {
148c:    83 f9 01                    cmp    $0x1,%ecx
148f:    7e 2c                        jle    14bd <do_loops+0x4d>
1491:    44 8d 41 fe                lea    -0x2(%rcx),%r8d
1495:    b8 01 00 00 00             mov    $0x1,%eax
149a:    49 83 c0 02                add    $0x2,%r8
149e:    66 90                        xchg   %ax,%ax
b[i] = a[i+1] + 3;
14a0:    8b 4c 87 04                mov    0x4(%rdi,%rax,4),%ecx
14a4:    83 c1 03                    add    $0x3,%ecx
14a7:    89 0c 86                    mov    %ecx,(%rsi,%rax,4)
c[i] = b[i-1] + 2;
14aa:    8b 4c 86 fc                mov    -0x4(%rsi,%rax,4),%ecx
14ae:    83 c1 02                    add    $0x2,%ecx
14b1:    89 0c 82                    mov    %ecx,(%rdx,%rax,4)
for (i=1; i<N; i++) {
14b4:    48 83 c0 01                add    $0x1,%rax
14b8:    4c 39 c0                    cmp    %r8,%rax
14bb:    75 e3                        jne    14a0 <do_loops+0x30>
}
14bd:    c3                          retq
14be:    66 90                        xchg   %ax,%ax
```

Analysis:

Compare to the original code with O2, O3, the O2 is a little bit faster than original one, while O3 is still slower than before.

When referring to assembly code: In the second and third loop (since the first loop doesn't change) in **original code**: we need to do the preparation work (from **cmp \$0x1, %ecx** to **xchg %ax, %ax**) twice (one for each loop initialization). For operation on `b[i]` and `c[i]`, original code use `%ecx` to store value for `b`, `%edi` to store value for `c`. And use **lea 0x2(%rdi), %ecx** to plus 2, which need an additional read from `%rdi`.

In **fusion code**, we only initialize once to access the address of `array[a]`, that we reduce overhead of loop management and improve data locality. Also we could reuse the `%ecx` to store value for `b` and `c`, thus cut down additional register usage.

Optimization 3: loop peeling

Source code:

```
int i;
if(N>=3) {
    a[N-1] = a[N-1] + 1;
}
for (i=N-2; i>=1; i--) {
    a[i] = a[i] + 1;
}
if(N>1) {
    b[1] = a[2] + 3;
}
for (i=2; i<N; i++) {
    b[i] = a[i+1] + 3;
}
if(N>1) {
    c[1] = b[0] + 2;
}
for (i=2; i<N; i++) {
    c[i] = b[i-1] + 2;
}
```

Performance result:

| Exe time(millisecons) | Optimization level | |
|-----------------------|--------------------|---------|
| | O2 | O3 |
| Cmd line argument | | |
| 10,000,000 | 22.629 | 18.912 |
| 100,000,000 | 225.357 | 187.003 |

Assembly Code of O2:

```
if(N>=3) {
1474:      8d 41 fe      lea    -0x2(%rcx),%eax
1477:      83 f9 02      cmp    $0x2,%ecx
147a:      0f 8e 88 00 00 00  jle    1508 <do_loops+0x98>
    a[N-1] = a[N-1] + 1;
1480:      4c 63 c1      movslq %ecx,%r8
1483:      42 83 44 87 fc 01  addl   $0x1,-0x4(%rdi,%r8,4)
for (i=N-2; i>=1; i--) {
1489:      48 98          cltq
148b:      0f 1f 44 00 00      nopl   0x0(%rax,%rax,1)
    a[i] = a[i] + 1;
1490:      83 04 87 01      addl   $0x1,(%rdi,%rax,4)
for (i=N-2; i>=1; i--) {
1494:      48 83 e8 01      sub    $0x1,%rax
1498:      85 c0          test   %eax,%eax
149a:      7f f4          jg     1490 <do_loops+0x20>
```

```

if(N>1) {
149c:      83 f9 01      cmp     $0x1,%ecx
149f:      7e 7f          jle     1520 <do_loops+0xb0>
    b[1] = a[2] + 3;
14a1:      8b 47 08      mov     0x8(%rdi),%eax
14a4:      83 c0 03      add     $0x3,%eax
14a7:      89 46 04      mov     %eax,0x4(%rsi)
    for (i=2; i<N; i++) {
14aa:      83 f9 02      cmp     $0x2,%ecx
14ad:      7e 71          jle     1520 <do_loops+0xb0>
14af:      44 8d 49 fd      lea     -0x3(%rcx),%r9d
14b3:      b8 03 00 00 00    mov     $0x3,%eax
14b8:      4d 89 c8      mov     %r9,%r8
14bb:      49 83 c1 04      add     $0x4,%r9
14bf:      90            nop
    b[i] = a[i+1] + 3;
14c0:      8b 0c 87      mov     (%rdi,%rax,4),%ecx
14c3:      83 c1 03      add     $0x3,%ecx
14c6:      89 4c 86 fc      mov     %ecx,-0x4(%rsi,%rax,4)

    for (i=2; i<N; i++) {
14ca:      48 83 c0 01      add     $0x1,%rax
14ce:      49 39 c1      cmp     %rax,%r9
14d1:      75 ed          jne     14c0 <do_loops+0x50>
    c[1] = b[0] + 2;
14d3:      8b 06      mov     (%rsi),%eax
14d5:      44 89 c1      mov     %r8d,%ecx
14d8:      48 8d 79 03      lea     0x3(%rcx),%rdi
14dc:      83 c0 02      add     $0x2,%eax
14df:      89 42 04      mov     %eax,0x4(%rdx)
14e2:      b8 02 00 00 00    mov     $0x2,%eax
14e7:      66 0f 1f 84 00 00 00    nopw    0x0(%rax,%rax,1)
14ee:      00 00
    c[i] = b[i-1] + 2;
14f0:      8b 4c 86 fc      mov     -0x4(%rsi,%rax,4),%ecx
14f4:      83 c1 02      add     $0x2,%ecx
14f7:      89 0c 82      mov     %ecx,(%rdx,%rax,4)
    for (i=2; i<N; i++) {
14fa:      48 83 c0 01      add     $0x1,%rax
14fe:      48 39 f8      cmp     %rdi,%rax
1501:      75 ed          jne     14f0 <do_loops+0x80>
1503:      c3            retq
1504:      0f 1f 40 00      nopl    0x0(%rax)

    for (i=N-2; i>=1; i--) {
1508:      85 c0      test    %eax,%eax
150a:      0f 8f 79 ff ff ff    jg      1489 <do_loops+0x19>
    if(N>1) {
1510:      83 f9 02      cmp     $0x2,%ecx
1513:      74 8c          je      14a1 <do_loops+0x31>
    }
1515:      c3            retq
1516:      66 2e 0f 1f 84 00 00    nopw    %cs:0x0(%rax,%rax,1)
151d:      00 00 00
    if(N>1) {
1520:      83 f9 02      cmp     $0x2,%ecx
1523:      75 f0          jne     1515 <do_loops+0xa5>
    c[1] = b[0] + 2;
1525:      8b 06      mov     (%rsi),%eax
1527:      83 c0 02      add     $0x2,%eax
152a:      89 42 04      mov     %eax,0x4(%rdx)
    for (i=2; i<N; i++) {
152d:      c3            retq
152e:      66 90      xchg    %ax,%ax

```

Analysis:

Compare to the original code with O2, O3, the loop peeling version is a little bit faster than the original ones in both O2 and O3 optimization(almost the same).

When referring to assembly code: the **original code** initialize the address at the beginning of each loop, then implement the add operation for corresponding operations.

In **loop peeling** version, however, 1. It enforce mem alignment for array before loop, 2. codes are interweaved. For example, we enter the first if sentence and finish the assignment,

then the loop for $a[i]$. Later, the second if sentence, and loop for $b[i]$. However, the `jump(loop judgment)` sentence for $b[i]$ operation resides in the next loop block, as a simplified way of `if()`. After the above section, there are code to check if we miss any circumstance due to this interweaving structure, if it does, jump back to previous section.

Optimization 4: loop reversal

Source code:

```
// loop reversal
int i;
for (i=1; i<N; i++) {
    a[i] = a[i] + 1;
}
for (i=N-1; i>=1; i--) {
    b[i] = a[i+1] + 3;
}
for (i=N-1; i>=1; i--) {
    c[i] = b[i-1] + 2;
}
```

Performance result:

| Exe time(milliseconds) | Optimization level | |
|------------------------|--------------------|---------|
| | O2 | O3 |
| Cmd line argument | | |
| 10,000,000 | 21.867 | 18.528 |
| 100,000,000 | 215.914 | 185.117 |

Assembly Code of O2:

```

1470:      f3 0f 1e fa      endbr64
for (i=1; i<N; i++) {
1474:      8d 41 ff          lea    -0x1(%rcx),%eax
1477:      83 f9 01          cmp    $0x1,%ecx
147a:      7e 5c              jle    14d8 <do_loops+0x68>
147c:      83 e9 02          sub    $0x2,%ecx
147f:      4c 8d 47 04          lea    0x4(%rdi),%r8
1483:      48 8d 4c 8f 08      lea    0x8(%rdi,%rcx,4),%rcx
1488:      0f 1f 84 00 00 00 00 nopl   0x0(%rax,%rax,1)
148f:      00
a[i] = a[i] + 1;
1490:      41 83 00 01          addl   $0x1, (%r8)
for (i=1; i<N; i++) {
1494:      49 83 c0 04          add    $0x4,%r8
1498:      4c 39 c1          cmp    %r8,%rcx
149b:      75 f3              jne    1490 <do_loops+0x20>
149d:      48 98              cltq
{
149f:      48 89 c1          mov    %rax,%rcx
14a2:      66 0f 1f 44 00 00      nopw   0x0(%rax,%rax,1)
b[i] = a[i+1] + 3;
14a8:      44 8b 4c 8f 04          mov    0x4(%rdi,%rcx,4),%r9d
14ad:      45 8d 41 03          lea    0x3(%r9),%r8d
14b1:      44 89 04 8e          mov    %r8d, (%rsi,%rcx,4)
for (i=N-1; i>=1; i--) {
14b5:      48 83 e9 01          sub    $0x1,%rcx
14b9:      85 c9              test   %ecx,%ecx
14bb:      7f eb              jg     14a8 <do_loops+0x38>
14bd:      0f 1f 00          nopl   (%rax)
c[i] = b[i-1] + 2;
14c0:      8b 7c 86 fc          mov    -0x4(%rsi,%rax,4),%edi
14c4:      8d 4f 02          lea    0x2(%rdi),%ecx
14c7:      89 0c 82          mov    %ecx, (%rdx,%rax,4)
for (i=N-1; i>=1; i--) {
14ca:      48 83 e8 01          sub    $0x1,%rax
14ce:      85 c0              test   %eax,%eax
14d0:      7f ee              jg     14c0 <do_loops+0x50>
}
14d2:      c3              retq
14d3:      0f 1f 44 00 00      nopl   0x0(%rax,%rax,1)
for (i=N-1; i>=1; i--) {
14d8:      85 c0              test   %eax,%eax
14da:      7f c1              jg     149d <do_loops+0x2d>
}
14dc:      c3              retq
14dd:      0f 1f 00          nopl   (%rax)

```

Analysis:

Compare to the original code with O2, O3, the loop reversal version is faster than the original ones in both O2 and O3 optimization.

In this case, the original code's assembly code is very similar to the reversal one, while the reversal one may enable more optimization choices to compiler.

Optimization 5: loop strip mining

Source code:

```
// loop strip mining
int i;
int j;
for (i=N-1; i>=1; i-=32) {
    for(j=i;j>(i-32);j--) {
        a[j] = a[j] + 1;
    }
}
for (i=1; i<N; i+=32) {
    for(j=i; j<(i+32); j++) {
        b[j] = a[j+1] + 3;
    }
}
for (i=1; i<N; i+=32) {
    for(j=i; j<(i+32); j++) {
        c[j] = b[j-1] + 2;
    }
}
```

Performance result:

| Exe time(millisecons) | Optimization level | |
|-----------------------|--------------------|---------|
| Cmd line argument | O2 | O3 |
| 10,000,000 | 24.191 | 17.921 |
| 100,000,000 | 236.418 | 182.449 |

Assembly Code of O2:

```
for (i=N-1; i>=1; i-=32) {
1477:      8d 49 ff          lea    -0x1(%rcx),%ecx
{
147a:      49 89 d0          mov    %rdx,%r8
for (i=N-1; i>=1; i-=32) {
147d:      85 c9             test   %ecx,%ecx
147f:      7e 47             jle    14c8 <do_loops+0x58>
1481:      48 63 c1          movslq %ecx,%rax
1484:      45 8d 4a df      lea    -0x21(%r10),%r9d
1488:      48 8d 54 87 80    lea    -0x80(%rdi,%rax,4),%rdx
148d:      41 8d 42 fe      lea    -0x2(%r10),%eax
1491:      83 e0 e0          and    $0xffffffff,%eax
1494:      41 29 c1          sub    %eax,%r9d
1497:      66 0f 1f 84 00 00 nopw   0x0(%rax,%rax,1)
149e:      00 00
for(j=i;j>(i-32);j--) {
14a0:      48 8d 82 80 00 00 lea     0x80(%rdx),%rax
14a7:      66 0f 1f 84 00 00 nopw   0x0(%rax,%rax,1)
14ae:      00 00
a[j] = a[j] + 1;
14b0:      83 00 01          addl   $0x1, (%rax)
for(j=i;j>(i-32);j--) {
14b3:      48 83 e8 04       sub    $0x4,%rax
14b7:      48 39 c2          cmp    %rax,%rdx
14ba:      75 f4             jne    14b0 <do_loops+0x40>
for (i=N-1; i>=1; i-=32) {
14bc:      83 e9 20          sub    $0x20,%ecx
14bf:      48 83 c2 80       add    $0xffffffffffff80,%rdx
14c3:      41 39 c9          cmp    %ecx,%r9d
14c6:      75 d8             jne    14a0 <do_loops+0x30>
```

```

for (i=1; i<N; i+=32) {
14c8:    b9 84 00 00 00    mov     $0x84,%ecx
14cd:    41 b9 01 00 00 00    mov     $0x1,%r9d
14d3:    41 83 fa 01        cmp     $0x1,%r10d
14d7:    7e 68              jle     1541 <do_loops+0xd1>
14d9:    0f 1f 80 00 00 00    nopl    0x0(%rax)
for(j=i; j<(i+32); j++) {
14e0:    48 8d 41 80        lea     -0x80(%rcx),%rax
14e4:    0f 1f 40 00        nopl    0x0(%rax)
b[j] = a[j+1] + 3;
14e8:    8b 54 07 04        mov     0x4(%rdi,%rax,1),%edx
14ec:    83 c2 03           add     $0x3,%edx
14ef:    89 14 06           mov     %edx,(%rsi,%rax,1)
for(j=i; j<(i+32); j++) {
14f2:    48 83 c0 04        add     $0x4,%rax
14f6:    48 39 c1           cmp     %rax,%rcx
14f9:    75 ed             jne     14e8 <do_loops+0x78>

for (i=1; i<N; i+=32) {
14fb:    41 83 c1 20        add     $0x20,%r9d
14ff:    48 83 e9 80        sub     $0xfffffffffffff80,%rcx
1503:    45 39 ca           cmp     %r9d,%r10d
1506:    7f d8             jg      14e0 <do_loops+0x70>
1508:    b9 84 00 00 00    mov     $0x84,%ecx
150d:    bf 01 00 00 00    mov     $0x1,%edi
1512:    66 0f 1f 44 00 00    nopw    0x0(%rax,%rax,1)
for(j=i; j<(i+32); j++) {
1518:    48 8d 41 80        lea     -0x80(%rcx),%rax
151c:    0f 1f 40 00        nopl    0x0(%rax)
c[j] = b[j-1] + 2;
1520:    8b 54 06 fc        mov     -0x4(%rsi,%rax,1),%edx
1524:    83 c2 02           add     $0x2,%edx
1527:    41 89 14 00        mov     %edx,(%r8,%rax,1)
for(j=i; j<(i+32); j++) {
152b:    48 83 c0 04        add     $0x4,%rax
152f:    48 39 c1           cmp     %rax,%rcx
1532:    75 ec             jne     1520 <do_loops+0xb0>
for (i=1; i<N; i+=32) {
1534:    83 c7 20        add     $0x20,%edi
1537:    48 83 e9 80        sub     $0xfffffffffffff80,%rcx
153b:    41 39 fa           cmp     %edi,%r10d
153e:    7f d8             jg      1518 <do_loops+0xa8>
1540:    c3              retq
1541:    c3              retq
1542:    66 2e 0f 1f 84 00 00    nopw    %cs:0x0(%rax,%rax,1)
1549:    00 00 00
154c:    0f 1f 40 00        nopl    0x0(%rax)

```

Analysis:

Compare to the original code with O2, O3, the O3 is a little bit faster than original one, while O2 is slower than before.

Since we access a block size of 32 each time, we improve the locality of this function. In **strip mining code**, Take $a[j] = a[j] + 1$ for example, the loop assembly section's code is less than original ones, which may help to improve performance. However, since there are more for() loop sentence so we need more **lea**, **nopw** and other command to initialize, so this may cause O2 to be slower than original ones.

Optimization 6: function inline

Source code:

```
struct timeval start_time, end_time;
gettimeofday(&start_time, NULL);
//do_loops(a, b, c, N);

for (i=N-1; i>=1; i--) {
    a[i] = a[i] + 1;
}
for (i=1; i<N; i++) {
    b[i] = a[i+1] + 3;
}
for (i=1; i<N; i++) {
    c[i] = b[i-1] + 2;
}
gettimeofday(&end_time, NULL);

for (i=0; i<N; i++) {
    sum += c[i];
}

double elapsed_us = calc_time(start_time, end_time);
double elapsed_ms = elapsed_us / 1000.0;
printf("sum=%llu\n", sum);
printf("Time=%f milliseconds\n", elapsed_ms);

return 0;
```

Performance result:

| Exe time(millisecons) | Optimization level | |
|-----------------------|--------------------|---------|
| | O2 | O3 |
| Cmd line argument | | |
| 10,000,000 | 22.987 | 18.427 |
| 100,000,000 | 224.757 | 185.468 |

Assembly Code of O2:

```
for (i=N-1; i>=1; i--) {
    11be: 41 8d 46 ff      lea    -0x1(%r14),%eax
    11c2: 85 c0            test   %eax,%eax
    11c4: 7e 16           jle    11dc <main+0xdc>
    11c6: 48 98           cltq
    11c8: 0f 1f 84 00 00 00 00 00  nopl   0x0(%rax,%rax,1)
    11cf: 00
    a[i] = a[i] + 1;
    11d0: 83 04 83 01      addl   $0x1, (%rbx,%rax,4)
    for (i=N-1; i>=1; i--) {
    11d4: 48 83 e8 01      sub    $0x1,%rax
    11d8: 85 c0            test   %eax,%eax
    11da: 7f f4           jg     11d0 <main+0xd0>
    }
}
```

```
for (i=1; i<N; i++) {
    11dc: 41 83 fe 01      cmp    $0x1,%r14d
    11e0: 0f 8e 1f 01 00 00 00 00  jle    1305 <main+0x205>
    11e6: 41 8d 4e fe      lea    -0x2(%r14),%ecx
    11ea: b8 02 00 00 00 00 00 00  mov     $0x2,%eax
    11ef: 48 8d 71 03      lea    0x3(%rcx),%rsi
    11f3: 0f 1f 44 00 00 00 00 00  nopl   0x0(%rax,%rax,1)
    b[i] = a[i+1] + 3;
    11f8: 8b 3c 83         mov     (%rbx,%rax,4),%edi
    11fb: 8d 57 03         lea     0x3(%rdi),%edx
    11fe: 41 89 54 84 fc   mov     %edx,-0x4(%r12,%rax,4)
    for (i=1; i<N; i++) {
    1203: 48 83 c0 01      add     $0x1,%rax
    1207: 48 39 c6         cmp     %rax,%rsi
    120a: 75 ec           jne     11f8 <main+0xf8>
    120c: 31 c0           xor     %eax,%eax
    120e: 66 90           xchg    %ax,%ax
    }
}
```

```

for (i=1; i<N; i++) {
  c[i] = b[i-1] + 2;
1210:    41 8b 34 84          mov     (%r12,%rax,4),%esi
1214:    8d 56 02          lea     0x2(%rsi),%edx
1217:    89 54 85 04          mov     %edx,0x4(%rbp,%rax,4)
  for (i=1; i<N; i++) {
121b:    48 89 c2          mov     %rax,%rdx
121e:    48 83 c0 01          add     $0x1,%rax
1222:    48 39 d1          cmp     %rdx,%rcx
1225:    75 e9          jne     1210 <main+0x110>
}

```

Analysis:

Compare to the original code with O2, O3, the loop reversal version is a little bit slower than/ almost the same as the original ones in both O2 and O3 optimization.

According to the assembly code, the inline version generate less assembly code than original one, it also avoid function call here. Since it doesn't do loop optimization, so the performance should be similar to original one.

(d)

Among all six methods I applied for this function, only three of them could defeat the compiler with corresponding Optimization level for sure:

Loop unrolling: defeat compiler with O2

Loop reversal: defeat compiler with O2

Loop strip mining: defeat compiler with O3

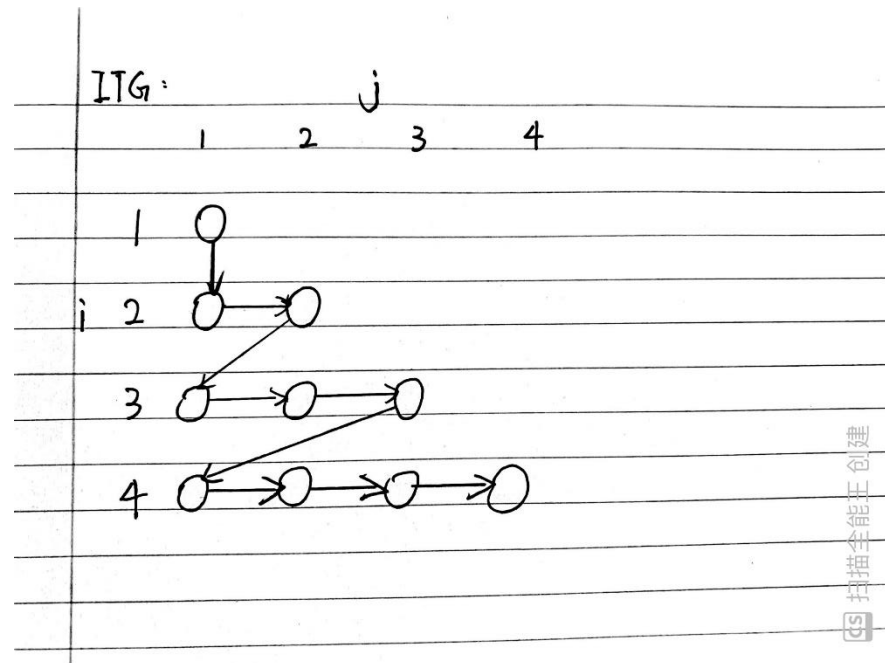
The performance of rest of methods are too close to the original code's, so it is hard to determine whether they defeat compiler for sure.

Question2 :

(a)

We assume $N = 4$ here

ITG of this code:



(b)

When $j=1, i=1$. Loop 1

S1: $a[1][1] = b[1][1] + c[1][1] * a[2][0];$

S2: $b[1][1] = a[0][0] * c[0][1];$

S3: $c[2][1] = a[1][1];$

S4: $d[1][1] = d[0][2];$

When $j=2, i=1$ Loop 2

S1: $a[2][1] = b[2][1] + c[2][1] * a[3][0];$

S2: $b[2][1] = a[1][0] * c[1][1];$

S3: $c[3][1] = a[2][1];$

S4: $d[2][1] = d[1][2];$

When $j=2, i=2$ Loop3

S1: $a[2][2] = b[2][2] + c[2][2] * a[3][1];$

S2: $b[2][2] = a[1][1] * c[1][2];$

S3: $c[3][2] = a[2][2]$;
S4: $d[2][2] = d[1][3]$;

When $j=3, i=1$ Loop4

S1: $a[3][1] = b[3][1] + c[3][1] * a[4][0]$;
S2: $b[3][1] = a[2][0] * c[2][1]$;
S3: $c[4][1] = a[3][1]$;
S4: $d[3][1] = d[2][2]$;

When $j=3, i=2$ Loop5

S1: $a[3][2] = b[3][2] + c[3][2] * a[4][1]$;
S2: $b[3][2] = a[2][1] * c[2][2]$;
S3: $c[4][2] = a[3][2]$;
S4: $d[3][2] = d[2][3]$;

Within the loop, we have loop-independent dependence:

S1 \rightarrow S2 (**anti**): $b[i][j]$
S1 \rightarrow S3 (**true**): $a[i][j]$


Between loop, we have loop-carried dependence

Loop1 S1 \rightarrow Loop3 S2 (**true**): $a[i][j]$ with $a[i-1][j-1]$;
Loop1 S3 \rightarrow Loop2 S1 (**true**): $c[i+1][j]$ with $c[i][j]$
Loop1 S3 \rightarrow Loop4 S2 (**true**): $c[i+1][j]$ with $c[i-1][j]$
Loop3 S4 \rightarrow Loop4 S4 (**true**): $d[i][j]$ with $d[i-1][j+1]$
Loop3 S1 \rightarrow Loop4 S1 (**anti**): $a[i+1][j-1]$ with $a[i][j]$

(c)

We still assume $N = 4$ here

LDG: 1 2 3 4



— : true dependence
 - - - : anti dependence

Question3 :

(a)

| Execute time(mili) | Mode for add() | | |
|--------------------|----------------|-----------|----------|
| Cmd line argument | inline | No-inline | original |
| 10,000,000 | 9.383 | 18.682 | 9.628 |
| 100,000,000 | 97.081 | 184.207 | 96.58 |

(b)

For inline version:

```
13e5: e8 e6 fd ff ff      callq 11d0 <gettimeofday@plt>
13ea: 48 8b 4c 24 50      mov 0x50(%rsp),%rcx
13ef: 48 8b 54 24 70      mov 0x70(%rsp),%rdx
13f4: 48 8b 74 24 30      mov 0x30(%rsp),%rsi
13f9: 48 8d 41 0f         lea 0xf(%rcx),%rax
13fd: 48 29 d0            sub %rdx,%rax
1400: 48 83 f8 1e         cmp $0x1e,%rax
1404: 48 8d 46 0f         lea 0xf(%rsi),%rax
1408: 40 0f 97 c7         seta %dil
140c: 48 29 d0            sub %rdx,%rax
140f: 48 83 f8 1e         cmp $0x1e,%rax
1413: 0f 97 c0            seta %al
1416: 40 84 c7            test %al,%dil
1419: 0f 84 72 02 00 00   je 1691 <main+0x451>
141f: 83 fd 02            cmp $0x2,%ebp
1422: 0f 86 69 02 00 00   jbe 1691 <main+0x451>
1428: 89 df              mov %ebx,%edi
142a: 31 c0              xor %eax,%eax
142c: c1 ef 02           shr $0x2,%edi
142f: 48 c1 e7 04        shl $0x4,%rdi
```

```

1433:    0f 1f 44 00 00    nopl    0x0(%rax,%rax,1)
1438:    f3 0f 6f 04 01    movdqu  (%rcx,%rax,1),%xmm0
143d:    f3 0f 6f 1c 06    movdqu  (%rsi,%rax,1),%xmm3
1442:    66 0f fe c3       paddb   %xmm3,%xmm0
1446:    0f 11 04 02       movups  %xmm0,(%rdx,%rax,1)
144a:    48 83 c0 10       add     $0x10,%rax
144e:    48 39 f8          cmp     %rdi,%rax
1451:    75 e5             jne     1438 <main+0x1f8>
1453:    89 d8             mov     %ebx,%eax
1455:    83 e0 fc          and     $0xffffffffc,%eax
1458:    f6 c3 03          test    $0x3,%bl
145b:    74 37             je      1494 <main+0x254>
145d:    48 63 f8          movslq  %eax,%rdi
1460:    44 8b 04 b9       mov     (%rcx,%rdi,4),%r8d
1464:    44 03 04 be       add     (%rsi,%rdi,4),%r8d
1468:    44 89 04 ba       mov     %r8d,(%rdx,%rdi,4)
146c:    8d 78 01          lea     0x1(%rax),%edi
146f:    39 fb             cmp     %edi,%ebx
1471:    7e 21             jle     1494 <main+0x254>
1473:    48 63 ff          movslq  %edi,%rdi
1476:    83 c0 02          add     $0x2,%eax
1479:    44 8b 04 be       mov     (%rsi,%rdi,4),%r8d
147d:    44 03 04 b9       add     (%rcx,%rdi,4),%r8d
1481:    44 89 04 ba       mov     %r8d,(%rdx,%rdi,4)
1485:    39 c3             cmp     %eax,%ebx
1487:    7e 0b             jle     1494 <main+0x254>
1489:    48 98             cltq
148b:    8b 34 86          mov     (%rsi,%rax,4),%esi
148e:    03 34 81          add     (%rcx,%rax,4),%esi
1491:    89 34 82          mov     %esi,(%rdx,%rax,4)
1494:    31 f6             xor     %esi,%esi
1496:    4c 89 e7          mov     %r12,%rdi
1499:    e8 32 fd ff ff    callq   11d0 <gettimeofday@plt>

```

For non_inline version:

```

13e8:    e8 e3 fd ff ff    callq   11d0 <gettimeofday@plt>
13ed:    4c 8b 4c 24 50     mov     0x50(%rsp),%r9
13f2:    4c 8b 44 24 30     mov     0x30(%rsp),%r8
13f7:    31 d2             xor     %edx,%edx
13f9:    48 8b 4c 24 70     mov     0x70(%rsp),%rcx
13fe:    66 90             xchgb   %ax,%ax
1400:    41 8b 34 91       mov     (%r9,%rdx,4),%esi
1404:    41 8b 3c 90       mov     (%r8,%rdx,4),%edi
1408:    e8 33 04 00 00     callq   1840 <_Z3addii>
140d:    89 04 91          mov     %eax,(%rcx,%rdx,4)
1410:    48 89 d0          mov     %rdx,%rax
1413:    48 83 c2 01       add     $0x1,%rdx
1417:    4c 39 e8          cmp     %r13,%rax
141a:    75 e4             jne     1400 <main+0x1c0>
141c:    31 f6             xor     %esi,%esi
141e:    4c 89 e7          mov     %r12,%rdi
1421:    e8 aa fd ff ff    callq   11d0 <gettimeofday@plt>

```

(c)

I expect the inline version runs much faster than no-inline version

According to the screen shoot above, it matches my expectation. Although the no-inline version seems shorter, it call add function by **callq 1840 <_Z3addii>** commend and jump there to execute. When it finish, jumps back. While in inline version, the “add” function is done “in place”, without additional execution jumping, saving a lot of time.

(d)

According to the performance form I record, the running time of original code is very close to the inline version, so compiler should use “in-lining” by default.

Question4 :

Original code:

```
Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    For (j=0; j<N; j++) {
        If (threshold < 4) {
            Sum = sum + a[j][i];
        } else {
            Sum = sum + a[j][i] + 1;
        }
    }
}
```

Opimization1: Loop Invariant hoisting

```
Int a[N][4];
Int rand_number = rand();
Threshold = 2.0 * rand_number;
For (i=0; i<4; i++) {
    For (j=0; j<N; j++) {
        If (threshold < 4) {
            Sum = sum + a[j][i];
```

```

        } else {
            Sum = sum + a[j][i] + 1;
        }
    }
}

```

Opimization2: Loop Unrolling

```

Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    If (threshold < 4) {
        Sum = sum + a[j][i];
        Sum = sum + a[j+1][i];
        ...
        Sum = sum + a[j+N-1][i];
    } else {
        Sum = sum + a[j][i] + 1;
        Sum = sum + a[j+1][i] + 1;
        ...
        Sum = sum + a[j+N-1][i] + 1;
    }
}
}

```

Opimization3: Loop Peeling

```

Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    If (threshold < 4) {
        Sum = sum + a[0][i];
    } else {
        Sum = sum + a[0][i] + 1;
    }
    For (j=1; j<N; j++) {
        If (threshold < 4) {
            Sum = sum + a[j][i];
        } else {
            Sum = sum + a[j][i] + 1;
        }
    }
}

```

```
}
```

Opimization4: Loop Unswitching

```
Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    If (threshold < 4) {
        For (j=0; j<N; j++) {
            Sum = sum + a[j][i];
        }
    } else {
        For (j=0; j<N; j++) {
            Sum = sum + a[j][i] + 1;
        }
    }
}
```

Opimization5: Loop Interchange

```
Int a[N][4];
Int rand_number = rand();
For (j=0; j<N; j++) {
    Threshold = 2.0 * rand_number;
    For (i=0; i<4; i++) {
        If (threshold < 4) {
            Sum = sum + a[j][i];
        } else {
            Sum = sum + a[j][i] + 1;
        }
    }
}
```

Opimization6: Loop Unroll and Jam

```
Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    For (j=0; j<N; j+=2) {
```

```

    If (threshold < 4) {
        Sum = sum + a[j][i];
        Sum = sum + a[j+1][i];
    } else {
        Sum = sum + a[j][i] + 1;
        Sum = sum + a[j+1][i] + 1;
    }
}
}

```

Opimization7: Loop Strip and Mining

```

Int a[N][4];
Int rand_number = rand();
For (i=0; i<4; i++) {
    Threshold = 2.0 * rand_number;
    For (j=0; j<N; j+=32) {
        For(int k = j; k < (j+32); k++) {
            If (threshold < 4) {
                Sum = sum + a[j][i];
            } else {
                Sum = sum + a[j][i] + 1;
            }
        }
    }
}
}

```


Question4 :

(a) Loop fusion:

Unsafe

In original code, there exist an anti-dependence between S2 and S3. After the transformation, the dependence changes to loop-carried output dependence between S3 and S1.

(b) Loop Interchange:

Unsafe

In original code, there exist an loop carried anti-dependence between S1, $i=1, j=2$ and S2, $i=2, j=1$. ($a[1][2] = a[2][1]$ and $a[2][1] = a[3][0]$). After the transformation, it changes to loop-carried true dependence between these two instruction ($j=1, i=2$ and $j=2, i=1$) ($a[2][1] = a[2][1]$ and $a[1][2] = a[2][1]$).

(c) Loop Fission:

Safe

In original code, it's not a dependence cycle. There is a loop carried true dependence between $i=1$ S1 with $i=3$ S2. After the transformation, it becomes loop-independence true dependence between S1 and S2.